



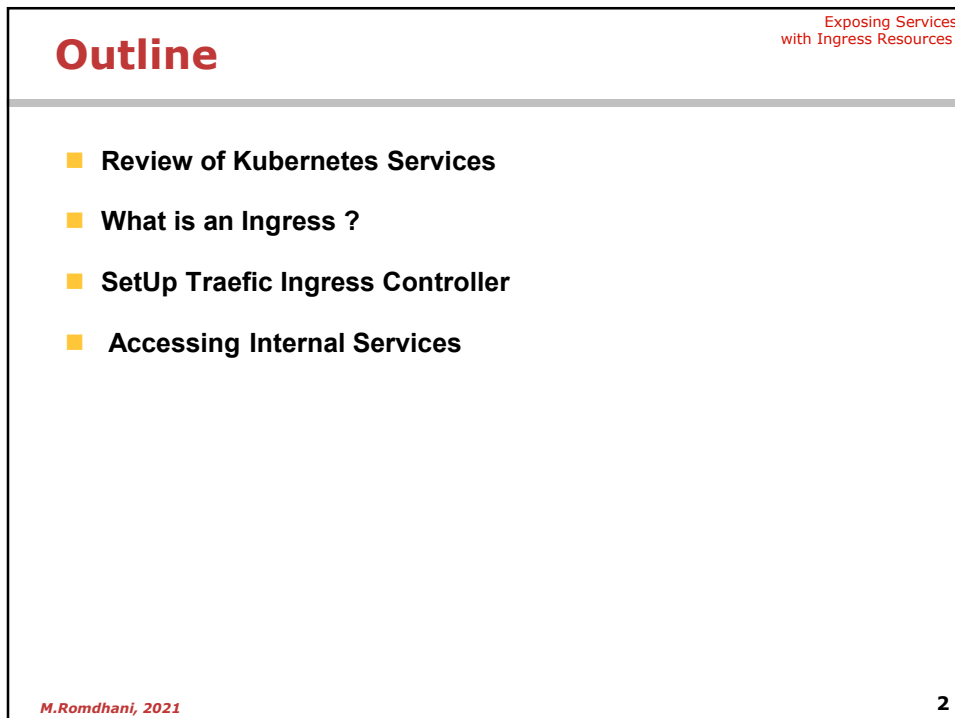
The slide features a purple header and footer bar. The header bar contains a small collage of images on the left. The main content area is white and contains a blue Kubernetes logo at the top center, followed by the text "Unit 5" in red. The title "Exposing Services with Ingress Resources" is written in a large, bold, red font in the center. In the bottom right corner, there are three small icons (a circle, a square, and a triangle) above the text "Business Training".

Unit 5

# Exposing Services with Ingress Resources

Business Training

1



The slide has a white background with a red header bar. The header bar contains the word "Outline" in a large, bold, red font on the left and the text "Exposing Services with Ingress Resources" in a smaller, red font on the right. Below the header bar, there is a list of four items, each preceded by a yellow square bullet point. At the bottom left, the text "M.Romdhani, 2021" is written in a small, red font. At the bottom right, the number "2" is written in a small, red font.

## Outline

Exposing Services with Ingress Resources

- Review of Kubernetes Services
- What is an Ingress ?
- SetUp Traffic Ingress Controller
- Accessing Internal Services

M.Romdhani, 2021

2

2

## Review of Kubernetes Services

3

### Services

Exposing Services  
with Ingress Resources

- **Services give us a [stable endpoint](#) to connect to a pod or a group of pods**
  - Durable resource (unlike Pods)
    - static cluster-unique IP
  - Target Pods using equality [based selectors](#)
  - kube-proxy provides simple load-balancing.
- **An easy way to create a service is to use `kubectl expose`**
  - If we have a deployment named my-little-deploy, we can run:  
`kubectl expose deployment my-little-deploy --port=80`
    - ... and this will create a service with the same name (my-little-deploy)
    - Services are automatically added to an internal DNS zone
    - In the example above, our code can now connect to <http://my-little-deploy/>
- **A service has a number of "endpoints"**

M.Romdhani, 2021

4

4

## Service Types

Exposing Services  
with Ingress Resources

- There are 3 major service types:
  1. ClusterIP (default)
  2. NodePort
  3. LoadBalancer
- There is also another resource type called Ingress (specifically for HTTP services)

M.Romdhani, 2021

5

5

## ClusterIP Services

Exposing Services  
with Ingress Resources

- It is the default service type
- A virtual IP address is allocated for the service
- This IP address is reachable only from within the cluster (nodes and pods)
- Perfect for internal communication, within the cluster

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: ClusterIP
  selector:
    app: nginx
    env: prod
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80 # should meet the
                    Container port
```

M.Romdhani, 2021

6

6

## NodePort Services

Exposing Services  
with Ingress Resources

- **NodePort services extend the ClusterIP service.**
  - Exposes a port on every node's IP.
- **Port can either be statically defined, or dynamically taken from a range between 30000-32767.**

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: NodePort
  selector:
    app: nginx
    env: prod
  ports:
    - nodePort: 30008
      port: 80
      targetPort: 80
```

M.Romdhani, 2021

7

7

## LoadBalancer Services

Exposing Services  
with Ingress Resources

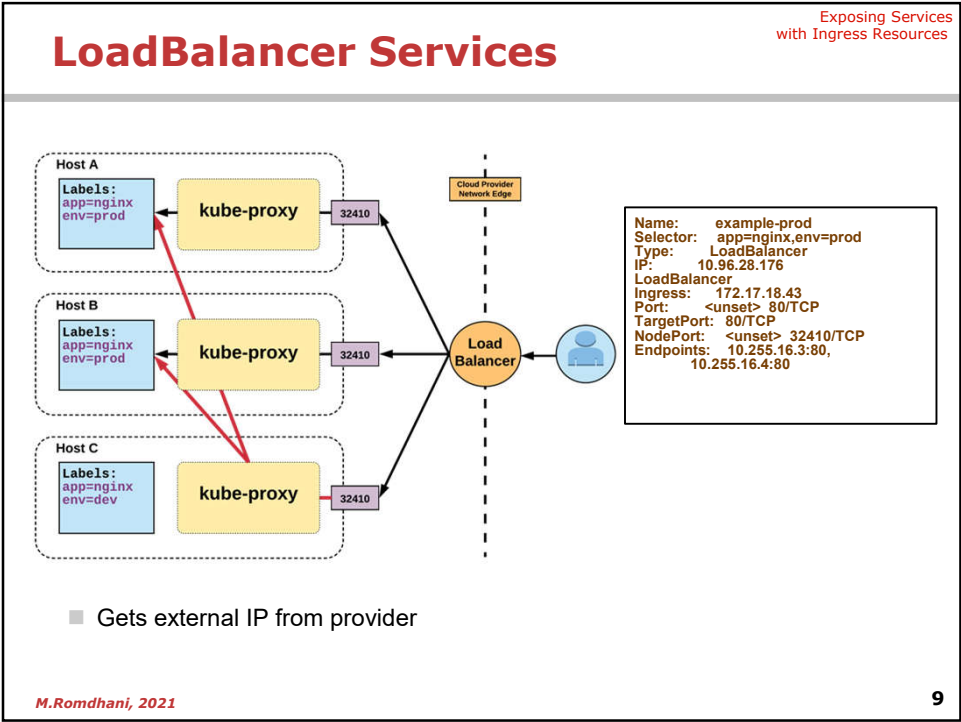
- **LoadBalancer services extend NodePort.**
- **Works in conjunction with an external system to map a cluster external IP to the exposed service (typically a cloud load balancer, e.g. ELB on AWS, GLB on GCE ...)**

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: LoadBalancer
  selector:
    app: nginx
    env: prod
  ports:
    protocol: TCP
    port: 80
    targetPort: 80
```

M.Romdhani, 2021

8

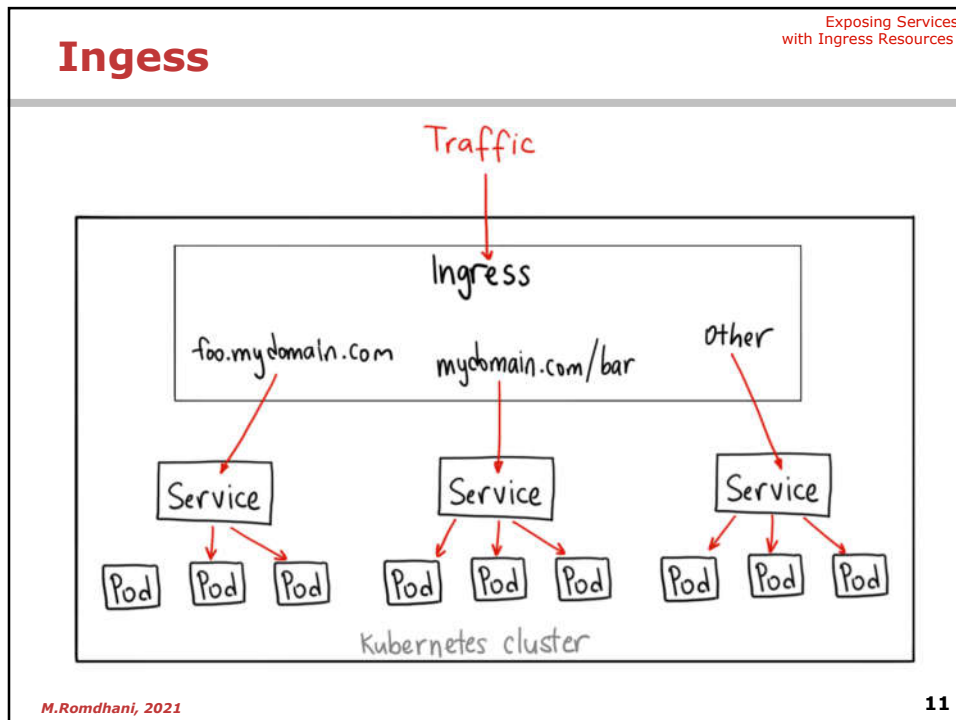
8



9

**What is an Ingress ?**

10



11

## The Ingress API Object

Exposing Services with Ingress Resources

- **An Ingress is an API object that manages external access to the services in a cluster**
  - Provides load balancing, SSL termination and name/path-based virtual hosting
  - Gives services externally-reachable URLs
- **They are specifically for HTTP services(not TCP or UDP)**
- **They can also handle TLS certificates, URL rewriting ...**

```
# Path based routing Example
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: simple-fanout-example
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - path: /foo
        backend:
          serviceName: service1
          servicePort: 4200
      - path: /bar
        backend:
          serviceName: service2
          servicePort: 8080
```

12

M.Romdhani, 2021

12

## Ingress Controller

Exposing Services  
with Ingress Resources

- The Ingress manifest doesn't actually do anything on its own; you must deploy **an Ingress Controller** into your cluster to watch for these declarations and act upon them.
- Ingress controllers are pods, just like any other application, so they're part of the cluster and can see other pods. They're built using reverse proxies that have been active in the market for years.
  - So, you have your choice of an **HAProxy**, **traefic**, **NGINX** Ingress Controller, and so on. The underlying proxy gives it Layer 7 routing and load balancing capabilities.
- Being inside the cluster themselves, Ingress Controllers are susceptible to the same walled-in jail as other Kubernetes pods.
  - You need to expose them to the outside via a Service with a type of either NodePort or LoadBalancer.
  - However, now you have a single endpoint that all traffic goes through: one Service connected to one Ingress Controller, which, in turn, is connected to many internal pods.
  - The controller, having the ability to inspect HTTP requests, directs a client to the correct pod based on characteristics it finds, such as the URL path or the domain name.

M.Romdhani, 2021

13

13

## Set up Traefic Ingress controller

14

## Traefik Ingress Controller

Exposing Services  
with Ingress Resources

### ■ Traefik is a modern HTTP reverse proxy and load balancer that makes deploying microservices easy.

- Traefik integrates with your existing infrastructure components (Docker, Swarm mode, Kubernetes, Amazon ECS, ...) and configures itself automatically and dynamically. Pointing Traefik at your orchestrator should be the only configuration step you need.



### ■ Features:

- Continuous update of configuration (no restarts),
- Support for multiple load balancing algorithms,
- Web UI, metrics export,
- Support for various protocols, REST API, canary releases and so on.
- The support for Let's Encrypt certificates right out of the box is another nice feature.

M.Romdhani, 2021

15

15

## Setup Traefik

Exposing Services  
with Ingress Resources

[<https://doc.traefik.io/traefik/v1.7/user-guide/kubernetes/>]

### ■ Prerequisites

- A working Kubernetes cluster. If you want to follow along with this guide, you should setup minikube on your machine, as it is the quickest way to get a local Kubernetes cluster setup for experimentation and development.
- Setup ingress as an add-on. It can be enabled by the following command:  
`minikube addons enable ingress`

### ■ Role Based Access Control configuration

- You will need to authorize Traefik to use the Kubernetes API

```
kubectl apply -f
https://raw.githubusercontent.com/traefik/traefik/v1.7/examples/k8s/traefik-rbac.yaml
```

### ■ Deploy Traefik using a Deployment or DaemonSet

- To deploy Traefik to your cluster start by submitting one of the YAML files to the cluster with kubectl:

```
kubectl apply -f
https://raw.githubusercontent.com/traefik/traefik/v1.7/examples/k8s/traefik-deployment.yaml
```

M.Romdhani, 2021

16

16



## Setup Traefik

Exposing Services  
with Ingress Resources

[<https://doc.traefik.io/traefik/v1.7/user-guide/kubernetes/>]

### ■ Check the pods

- Start by listing the pods in the kube-system namespace:

```
kubectl --namespace=kube-system get pods
```

### ■ Submitting an Ingress to the Cluster

- Lets start by creating a Service and an Ingress that will expose the Traefik Web UI.

```
kubectl apply -f  
https://raw.githubusercontent.com/traefik/traefik/v1.7/examples/k8s/ui.yaml
```

In production you would want to set up real DNS entries. You can get the IP address of your minikube instance by running minikube ip:

```
echo "$(minikube ip) traefik-ui.minikube" | sudo tee -a /etc/hosts
```

- We should now be able to visit traefik-ui.minikube in the browser and view the Traefik web UI : <https://traefik-ui.minikube/>

M.Romdhani, 2021

17

17

## Accessing Internal Services

18

## Accessing internal services

Exposing Services  
with Ingress Resources

- **When we are logged in on a cluster node, we can access internal services**
  - As per the Kubernetes network model: all nodes can reach all pods and services)
- **When we are accessing a remote cluster, our local machine won't have access to the cluster's internal subnet. To overcome this:**
  - **kubectrl proxy**: gives us access to the API, which includes a proxy for HTTP resources
  - **kubectrl port-forward**: allows forwarding of TCP ports to arbitrary pods, services, ...

M.Romdhani, 2021

19

19

## kubectrl proxy

Exposing Services  
with Ingress Resources

- **Running **kubectrl proxy** gives us access to the entire Kubernetes API**
  - The API includes routes to proxy HTTP traffic
  - By default, the proxy listens on port 8001
- **These routes look like the following:**
  - `/api/v1/namespaces/<namespace>/services/<service>/proxy`
- **We just add the URI to the end of the request, for instance:**
  - `/api/v1/namespaces/<namespace>/services/<service>/proxy/index.html`
- **We can access services and pods this way !**
- **Security considerations : *kubectrl proxy is intended for local use***
  - Running kubectrl proxy openly is a huge security risk
  - It is slightly better to run the proxy where you need it (and copy credentials, e.g. ~/.kube/config, to that place)
  - It is even better to use a limited account with reduced permissions

M.Romdhani, 2021

20

20

## kubectrl port-forward

Exposing Services  
with Ingress Resources

### ■ What if we want to access a TCP service?

- We can use **kubectrl port-forward** instead
- It will create a TCP relay to forward connections to a specific port (of a pod, service, deployment...)

### ■ The syntax is:

`kubectrl port-forward service/name_of_service local_port:remote_port`

- If only one port number is specified, it is used for both local and remote ports

M.Romdhani, 2021

21